

C1_W5_Lab_2_custom-callbacks

March 16, 2021

1 Ungraded Lab: Keras custom callbacks

A custom callback is a powerful tool to customize the behavior of a Keras model during training, evaluation, or inference. Towards the end of this guide, there will be demos of creating a couple of simple callback applications to get you started on your custom callback.

1.1 Imports

```
[ ]: import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np

import datetime
import io

from PIL import Image
from IPython.display import Image as IPyImage
import imageio

print("Version: ", tf.__version__)
tf.get_logger().setLevel('INFO')
```

```
[ ]: # Define the Keras model to add callbacks to
def get_model():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Dense(1, activation = 'linear', input_dim = 784))
    model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=0.1),
↳ loss='mean_squared_error', metrics=['mae'])
    return model
```

Then, load the MNIST data for training and testing from Keras datasets API:

```
[ ]: # Load example MNIST data and pre-process it
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train = x_train.reshape(60000, 784).astype('float32') / 255
x_test = x_test.reshape(10000, 784).astype('float32') / 255
```

Now, define a simple custom callback to track the start and end of every batch of data. During those calls, it prints the index of the current batch.

```
[ ]: class MyCustomCallback(tf.keras.callbacks.Callback):

    def on_train_batch_begin(self, batch, logs=None):
        print('Training: batch {} begins at {}'.format(batch, datetime.datetime.
→now().time()))

    def on_train_batch_end(self, batch, logs=None):
        print('Training: batch {} ends at {}'.format(batch, datetime.datetime.
→now().time()))
```

Providing a callback to model methods such as `tf.keras.Model.fit()` ensures the methods are called at those stages:

```
[ ]: model = get_model()
_ = model.fit(x_train, y_train,
             batch_size=64,
             epochs=1,
             steps_per_epoch=5,
             verbose=0,
             callbacks=[MyCustomCallback()])
```

1.2 An overview of callback methods

1.2.1 Common methods for training/testing/predicting

For training, testing, and predicting, following methods are provided to be overridden. ##### `on_(train|test|predict)_begin(self, logs=None)` Called at the beginning of fit/evaluate/predict. ##### `on_(train|test|predict)_end(self, logs=None)` Called at the end of fit/evaluate/predict. ##### `on_(train|test|predict)_batch_begin(self, batch, logs=None)` Called right before processing a batch during training/testing/predicting. Within this method, `logs` is a dict with `batch` and `size` available keys, representing the current batch number and the size of the batch. ##### `on_(train|test|predict)_batch_end(self, batch, logs=None)` Called at the end of training/testing/predicting a batch. Within this method, `logs` is a dict containing the stateful metrics result.

1.2.2 Training specific methods

In addition, for training, following are provided. ##### `on_epoch_begin(self, epoch, logs=None)` Called at the beginning of an epoch during training. ##### `on_epoch_end(self, epoch, logs=None)` Called at the end of an epoch during training.

1.2.3 Usage of logs dict

The logs dict contains the loss value, and all the metrics at the end of a batch or epoch. Example includes the loss and mean absolute error.

```
[ ]: callback = tf.keras.callbacks.LambdaCallback(
    on_epoch_end=lambda epoch,logs:
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch,
            ↪logs["val_loss"] / logs["loss"]))
)

model = get_model()
_ = model.fit(x_train, y_train,
              validation_data=(x_test, y_test),
              batch_size=64,
              epochs=3,
              verbose=0,
              callbacks=[callback])

[ ]: class DetectOverfittingCallback(tf.keras.callbacks.Callback):
    def __init__(self, threshold=0.7):
        super(DetectOverfittingCallback, self).__init__()
        self.threshold = threshold

    def on_epoch_end(self, epoch, logs=None):
        ratio = logs["val_loss"] / logs["loss"]
        print("Epoch: {}, Val/Train loss ratio: {:.2f}".format(epoch, ratio))

        if ratio > self.threshold:
            print("Stopping training...")
            self.model.stop_training = True

model = get_model()
_ = model.fit(x_train, y_train,
              validation_data=(x_test, y_test),
              batch_size=64,
              epochs=3,
              verbose=0,
              callbacks=[DetectOverfittingCallback()])
```

Similarly, one can provide callbacks in `evaluate()` calls.

1.3 Custom callback to Visualize predictions

```
[ ]: # Load example MNIST data and pre-process it
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()
x_train = x_train.reshape(60000, 784).astype('float32') / 255
x_test = x_test.reshape(10000, 784).astype('float32') / 255

[ ]: # Visualization utilities
plt.rc('font', size=20)
plt.rc('figure', figsize=(15, 3))

def display_digits(inputs, outputs, ground_truth, epoch, n=10):
    plt.clf()

    plt.yticks([])
    plt.grid(None)
    inputs = np.reshape(inputs, [n, 28, 28])
    inputs = np.swapaxes(inputs, 0, 1)
    inputs = np.reshape(inputs, [28, 28*n])
    plt.imshow(inputs)
    plt.xticks([28*x+14 for x in range(n)], outputs)
    for i, t in enumerate(plt.gca().xaxis.get_ticklabels()):
        if outputs[i] == ground_truth[i]:
            t.set_color('green')
        else:
            t.set_color('red')
    plt.grid(None)

[ ]: GIF_PATH = './animation.gif'

[ ]: class VisCallback(tf.keras.callbacks.Callback):
    def __init__(self, inputs, ground_truth, display_freq=10, n_samples=10):
        self.inputs = inputs
        self.ground_truth = ground_truth
        self.images = []
        self.display_freq = display_freq
        self.n_samples = n_samples

    def on_epoch_end(self, epoch, logs=None):
        # Randomly sample data
        indexes = np.random.choice(len(self.inputs), size=self.n_samples)
        X_test, y_test = self.inputs[indexes], self.ground_truth[indexes]
        predictions = np.argmax(self.model.predict(X_test), axis=1)

        # Plot the digits
        display_digits(X_test, predictions, y_test, epoch, n=self.display_freq)
```

```

    # Save the figure
    buf = io.BytesIO()
    plt.savefig(buf, format='png')
    buf.seek(0)
    image = Image.open(buf)
    self.images.append(np.array(image))

    # Display the digits every 'display_freq' number of epochs
    if epoch % self.display_freq == 0:
        plt.show()

    def on_train_end(self, logs=None):
        imageio.mimsave(GIF_PATH, self.images, fps=1)

```

```

[ ]: def get_model():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Dense(32, activation='linear', input_dim=784))
    model.add(tf.keras.layers.Dense(10, activation='softmax'))
    model.compile(optimizer=tf.keras.optimizers.RMSprop(lr=1e-4),
    ↪ loss='sparse_categorical_crossentropy', metrics=['accuracy'])
    return model

```

```

[ ]: model = get_model()
model.fit(x_train, y_train,
        batch_size=64,
        epochs=20,
        verbose=0,
        callbacks=[VisCallback(x_test, y_test)])

```

```

[ ]: SCALE = 60

# FYI, the format is set to PNG here to bypass checks for acceptable embeddings
IPyImage(GIF_PATH, format='png', width=15 * SCALE, height=3 * SCALE)

```